

Not Kevin Bacon

Suppose we have a graph or network with n nodes and some connections between them. The matrix \mathbf{A} whose entries are

$$A_{ij} = \begin{cases} 1, & \text{if node } i \text{ connects to } j, \\ 0, & \text{otherwise,} \end{cases} \quad (1)$$

is the **adjacency matrix** of the network. Note that these connections have direction, and i connecting to j does not automatically imply that j connects to i (i.e., this is a directed graph). Let $s_i = \sum_{j=1}^n A_{ij}$ be the number of nodes that node i connects to. If we were to randomly select a link leaving node i , each link would have probability $1/s_i$ of being selected.

Let \mathbf{x} be a vector of positive values. We will require that $\sum_i x_i = 1$, so that \mathbf{x} has the interpretation of a probability distribution over the nodes. If all connections are given equal weighting, the probability of following a connection from any node to node i is

$$z_i = \sum_{j \in P_i} \frac{x_j}{s_j}, \quad (2)$$

where P_i is the set of nodes that connect to node i . These are the rows with ones in column i of \mathbf{A} . By the definition of \mathbf{A} , this is the same as

$$z_i = \sum_{j=1}^n \frac{A_{ji}x_j}{s_j} = \sum_{j=1}^n B_{ij}x_j, \quad (3)$$

where we defined $B_{ij} = A_{ji}/s_j$. Put simply, $\mathbf{z} = \mathbf{B}\mathbf{x}$.

It's important to introduce some overall randomness into the jumps between nodes—this is the only way to escape a self-contained clique (disconnected subgraph). The probability of hopping to any node i entirely at random is just $1/n$. We blend link-following with random hopping as follows. Choose some $p \in [0, 1]$, and suppose that a hop between nodes follows one of the connections with probability p , or is a random hop with probability $1 - p$. Using $\mathbf{1}$ to denote the n -vector of all ones, then

$$\mathbf{y} = p\mathbf{B}\mathbf{x} + \frac{1-p}{n}\mathbf{1} \quad (4)$$

describes how to update probabilities after each hop. Finally, the fact that

$$1 = \sum_i x_i = \mathbf{1}^T \mathbf{x}, \quad (5)$$

allows us to express the map (4) as

$$\mathbf{y} = \left[p\mathbf{B} + \frac{1-p}{n}\mathbf{1}\mathbf{1}^T \right] \mathbf{x} = \mathbf{R}\mathbf{x}, \quad (6)$$

for a square matrix \mathbf{R} . If the probabilities are unchanged by hopping (i.e., $\mathbf{z} = \mathbf{x}$), then \mathbf{x} is an eigenvector of \mathbf{R} with associated eigenvalue $\lambda = 1$. We won't prove this, but \mathbf{R} is guaranteed to have $\lambda = 1$ as the leading eigenvalue, making power iteration possible. The resulting eigenvector \mathbf{x} can be sorted to find out which nodes are most likely to be visited in the long run.

Note that \mathbf{R} is *not* sparse and should never be formed. However, $\mathbf{R}\mathbf{x}$ as defined in (4) can be computed efficiently if \mathbf{B} is sparse. That is all we need to do a power iteration with \mathbf{R} . Since we are working with probability, normalization is not required in the power iteration: $\mathbf{x}_{k+1} = \mathbf{R}\mathbf{x}_k$, provided \mathbf{x}_1 has positive entries and $\|\mathbf{x}\|_1 = 1$.

Goals

You will use an adjacency matrix for movie actors to perform the power iteration and find the leading eigenvector of \mathbf{R} , and using that vector to rank the actors in influence.

Preparation

Read section 8.2. Answer the following questions based on the above description.

1. Show using (3) that $\sum_{i=1}^n z_i = 1$. This proves that \mathbf{z} is also a probability distribution.
2. Show using (4) that \mathbf{y} is a probability distribution.

Procedure

Download the script template and the data file `actornetwork.mat`.

1. Load `actornetwork.mat` file from the assignment site. It has a vector `actor` of unique actor names for all credited roles in films released from 2004 through 2013. It also has a sparse adjacency matrix `A`, where a (symmetric) link between actors means that they appeared in at least one film together.

2. Use `nnz` to compute the density (number of nonzeros over total number of elements) of \mathbf{A} . Use `whos` to find the memory usage of \mathbf{A} in bytes. Also calculate the memory usage of an equivalent full (non-sparse) matrix.
3. Compute the vector \mathbf{s} whose entries are s_i as defined above. Make a histogram of its entries using 32 bins.
4. Construct the matrix \mathbf{B} appearing in (3) above. It helps to use the fact that \mathbf{A} is symmetric. (It's reasonable to loop over one dimension of the matrix, but not over all of the elements.)
5. Set $p = 0.9$ and let \mathbf{x}_1 be a random vector of positive numbers. Normalize \mathbf{x}_1 to be a probability distribution. By repeatedly applying (4), do 100 power iterations to get \mathbf{x}_{101} . *Important: Do not attempt to define the matrix \mathbf{R} , and do not use the book's power iteration function; instead use `(?)` to compute $\mathbf{R}\mathbf{x}$. Check that $\|\mathbf{x}_{101} - \mathbf{x}_{100}\|_1$ is less than 10^{-6} .*
6. Sort the entries of \mathbf{x} in descending order, using the second output to print out the names of the 10 “most collaborative” actors.

Discussion

1. The data file also includes an $m \times n$ sparse matrix \mathbf{M} . Its (i, j) entry is one if actor j appeared in film i , and zero otherwise. Give a simple interpretation of the matrix $\mathbf{M}^T\mathbf{M}$.
2. What is the interpretation of $\mathbf{M}\mathbf{M}^T$? What would a ranking using this matrix reveal?