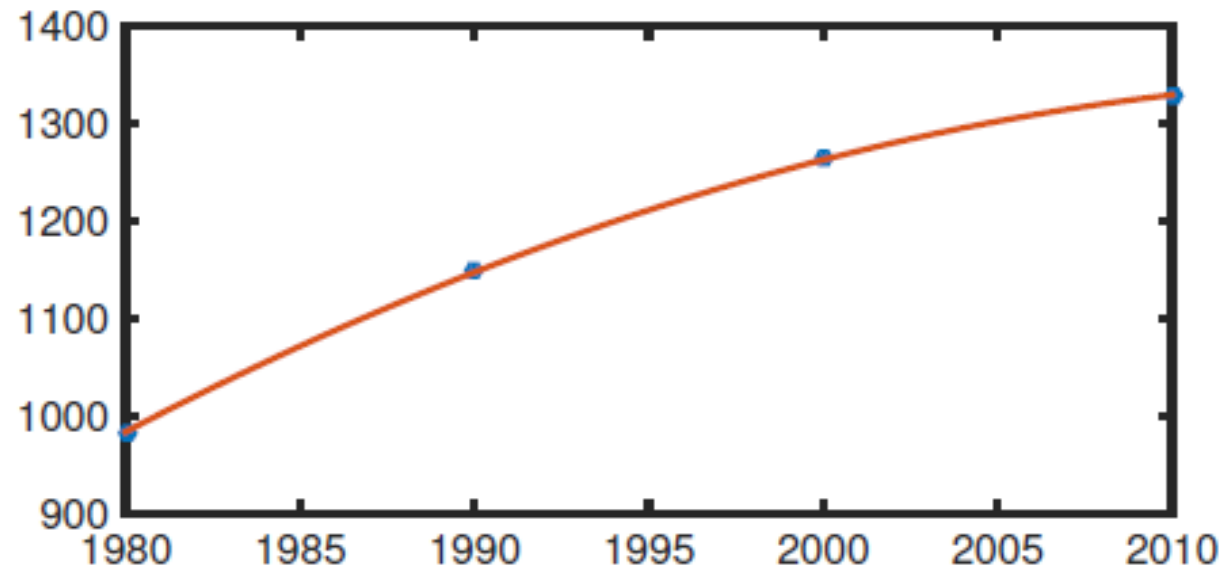


Chapter 2

Square linear systems: $Ax = b$



Comparing function sizes in a limit

- We want to compare how functions behave in limiting cases
- Let $f(n) = \tan(1/n)$, $g(n) = 1/n$ and $h(n) = 1/n^2$
- How do these compare to each other as $n \rightarrow \infty$, or “large n ”
- The limit is easy: $\lim_{n \rightarrow \infty} f(n) = \tan(0) = 0$
- What we want to know is how fast this function gets to the limit. We can do that by comparing f to g and h
- With g , $\lim_{n \rightarrow \infty} f(n)/g(n) = \lim_{n \rightarrow \infty} \tan\left(\frac{1}{n}\right) / \left(\frac{1}{n}\right) = 1$
- We conclude that f and g are the same size...

Comparing functions: Order and asymptotic

- We say that “ f is asymptotic to g ” or $f \sim g$
- These two approach the limit at the same rate
- We can also say “ f is order g ” or $f = O(g)$ when the limit is bounded
- With h , $\lim_{n \rightarrow \infty} \frac{f(n)}{h(n)} = \lim_{n \rightarrow \infty} [\tan\left(\frac{1}{n}\right)] / \left(\frac{1}{n^2}\right) = \lim_{n \rightarrow \infty} n = \infty$
- The limit doesn't exist, because h goes to zero faster than f
- So, f and h are not the same order.
- If we define $k(n) = \tan^2\left(\frac{1}{n}\right)$, we find $k \sim h$

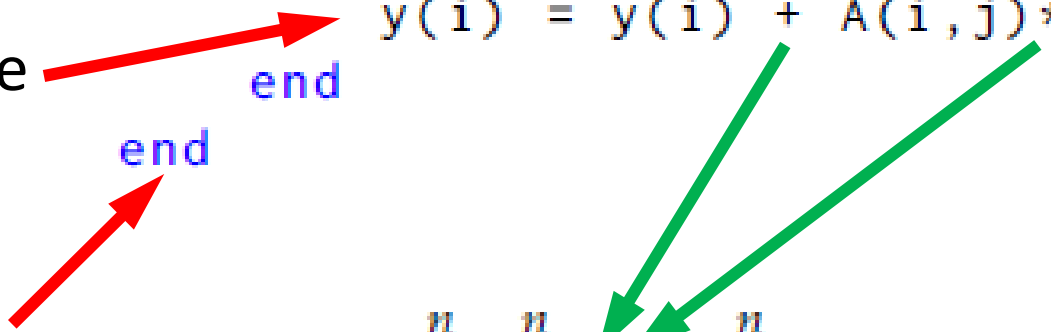
Comparing functions: Order and asymptotic

- The dominant part of a growing function can be identified using this approach
- Let $f(n) = a_1n^2 + b_1n + c_1$ and $g(n) = n^2$
- Then $\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = \lim_{n \rightarrow \infty} a_1 + b_1n^{-1} + c_1n^{-2} = a_1$
- We say that “ f is asymptotic to n^2 ” or $f \sim n^2$
- We’ll use this to evaluate operation counts and performance

Operation Counts: Example

- Matlab example multiplies a matrix and a vector, then adds a vector
- Count *,/,+,- as same operation
- Neglect storage
- Inner loop: Multiplying one row of A with x takes $n *$ and $n-1 +$; adding vector adds $1 +$
- Outer loop: adding that up for n rows gives final result of $2n^2$

```
n = 6;  
A = magic(n);  
x = ones(n,1);  
y = zeros(n,1);  
for i = 1:n  
    for j = 1:n  
        y(i) = y(i) + A(i,j)*x(j);  
    end  
end
```

$$\sum_{i=1}^n \sum_{j=1}^n 2 = \sum_{i=1}^n 2n = 2n^2.$$


Operation Counts: Matlab

- Let's test Matlab matrix-vector multiplication
- Time using stopwatch functions tic and toc
- Repeat to get more reliable timing

```
fprintf('    n    time (sec)\n')  
fprintf('    \n')  
fprintf('\n\n')
```

```
t_ = [];  
n_ = 400:400:4000;  
for n = n_  
    A = randn(n,n);    x = randn(n,1);  
    tic  
    for j = 1:10  
        A*x;  
    end  
    t = toc;  
    t_ = [t_, t/10];  
end
```

n	time (sec)
400	1.47e-04
800	7.49e-04
1200	1.80e-03
1600	1.88e-03
2000	2.69e-03
2400	4.14e-03
2800	5.29e-03
3200	4.81e-03
3600	7.75e-03
4000	7.40e-03

How fast is the time growing?


Operation Counts: Matlab

- Matlab matrix-vector multiplication
- How to identify rate? Expecting power law: use log-log plot.

$$t = Cn^p \implies \log t = p(\log n) + (\log C).$$

- If we get a line, the slope is p


```
loglog(n_,t_,'.-')  
xlabel('size of matrix'), ylabel('time (sec)')  
title('Timing of matrix-vector multiplications')
```



Operation Counts: Matlab

- To see the slope one could plot a known function
- Try Cn^2 here because of theory
- Pick C to put line in a convenient place, or use t_1/n_1^2

```
hold on, loglog(n_,t_(1)*(n_/n_(1)).^2,'--')  
axis tight  
legend('data','O(n^2)','location','southeast')
```



Operation Count: LU

- While some time is needed, no FLOPs here, so neglected

- Inside nested loops, one FLOP here

- Multiply part of row (j to n, for n-j+1 elements) and subtract from same part

- For operations inside both loops total is then

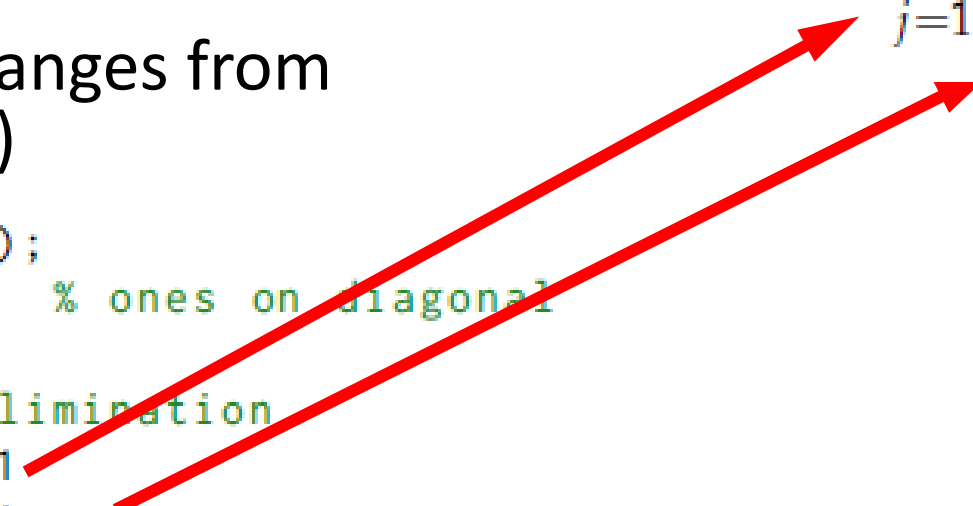
$$1 + 2(n - j + 1) \\ = 2(n - j) + 3$$

- Line 19 ignored
- Now total over loops

```
8  n = length(A);
9  L = eye(n);    % ones on diagonal
10
11 % Gaussian elimination
12 for j = 1:n-1
13     for i = j+1:n
14         L(i,j) = A(i,j) / A(j,j);    % row multiplier
15         A(i,j:n) = A(i,j:n) - L(i,j)*A(j,j:n);
16     end
17 end
18
19 U = triu(A);
```

Operation Count: LU

- Inner loop: i ranges from $j + 1$ to n (rows below diag)
- Outer loop: j ranges from 1 to n (all rows)

$$\sum_{j=1}^{n-1} \sum_{i=j+1}^n [2(n-j) + 3]$$


```
8  n = length(A);
9  L = eye(n);    % ones on diagonal
10
11 % Gaussian elimination
12 for j = 1:n-1
13     for i = j+1:n
14         L(i,j) = A(i,j) / A(j,j);    % row multiplier
15         A(i,j:n) = A(i,j:n) - L(i,j)*A(j,j:n);
16     end
17 end
18
19 U = triu(A);
```

Note how the index for the outer sum j is in the limit of the inner sum: We must do the inner sum first to get all the j 's into the summand

Operation Count: LU

- Do the inner sum:

$$\sum_{j=1}^{n-1} \sum_{i=j+1}^n [2(n-j) + 3] = \sum_{j=1}^{n-1} (n-j) [2(n-j) + 3]$$

- Make it easier: change variable with $k = n-j$.
- When $j = 1$, $k = n - 1$; when $j = n - 1$, $k = 1$. Then we have

$$\sum_{k=1}^{n-1} k(2k + 3)$$

- Distributing the sum, we have sum involving k and a sum involving k^2

Operation Count: LU

- For most situations, we only care about the leading factor in the sum, and can use the results at right

$$\sum_{k=1}^{n-1} k(2k+3)$$

- One term is proportional to n^3 , from k^2
- The other is proportional to n^2 , from k
- Our sum then results in two parts:

$$\frac{2n^3}{3} + \frac{3n^2}{2}$$

- For large n , the cubic term is dominant: LU factorization is $O(n^3)$, and more

$$\sum_{k=1}^n k \sim \frac{n^2}{2} = O(n^2), \text{ as } n \rightarrow \infty,$$

$$\sum_{k=1}^n k^2 \sim \frac{n^3}{3} = O(n^3), \text{ as } n \rightarrow \infty,$$

\vdots

$$\sum_{k=1}^n k^p \sim \frac{n^{p+1}}{p+1} = O(n^{p+1}), \text{ as } n \rightarrow \infty,$$

Operation Count: LU

- Our sum then results in two parts:

$$\frac{2n^3}{3} + \frac{3n^2}{2}$$

- For large n , the cubic term is dominant: LU factorization is $O(n^3)$
- More specifically, the FLOP count is asymptotic to
$$\frac{2n^3}{3}$$
- How does this stack up against actual computation time?

Operation Counts: LU

- Let's test with Matlab
- Use functions tic and toc: wall clock time
- Increase matrix size n
- Repeat to get more reliable timing


```
t_ = [];  
n_ = 200:100:2400;  
for n = n_  
    A = randn(n,n);  
    tic  
    for j = 1:6, [L,U] = lu(A); end  
    t = toc;  
    t_ = [t_, t/6];  
end
```



How fast is the time growing with n ?

Operation Counts: Matlab


```
loglog(n_,t_,'.-')
hold on, loglog(n_,t_(end)*(n_/n_(end)).^3,'--')
axis tight
xlabel('size of matrix'), ylabel('time (sec)')
title('Timing of LU factorization')
legend('lu','O(n^3)','location','southeast')
```



- Note that we picked a convenient constant
- Is this behaving like n^3 ?
- It's not perfect fit by any means

Operation Counts: Matlab

```
hold on, loglog(n_, t_(end)*(n_/n_(end)).^2, '--')  
legend('lu', 'O(n^3)', 'O(n^2)', 'location', 'southeast')
```



- Let's add an n^2 curve
- It looks like it is between n^2 and n^3
- For larger matrix size, seems closer to n^3
- What is contributing to this?

Operation Count: LU

- We assumed all that mattered was time to do one flop and that they were sequential
- CPUs can have multiple cores, and can have vectorized operations
- These can violate our assumptions
- Parallel computation is even more different: time to send data to different CPUs can even dominate the computation
- The details of the specific hardware matter
- If computational time is important to your project, test it!!!!

Fixing up issues with naïve GE

- To carry out GE, we needed to compute the multiplier
$$L(i + 1, i) = A(i + 1, i) / A(i, i).$$
- What if $A(i, i) = 0$?
- Switch rows so that there is nonzero element there:
“row pivoting”
- Smart way to do that: move the largest element of $A(i+1:n, i)$, the part of column i below the pivot, to the pivot.
- Consider this example...

Our previous example

```
A = [2 0 4 3 ; -4 5 -7 -10 ; 1 15 2 -4.5 ; -2 0 2 -13]
b = [ 4; 9; 29; 40 ]
```

```
A =
    2.0000         0    4.0000    3.0000
   -4.0000    5.0000   -7.0000  -10.0000
    1.0000   15.0000    2.0000   -4.5000
   -2.0000         0    2.0000  -13.0000

b =
     4
     9
    29
    40
```

```
[L,U] = lufact(A);
x = backsub( U, forwardsub(L,b) )
```

```
x =
    -3
     1
     4
    -2
```

- No difficulties here
- But, after finishing with the first column, we had $A_{24} = 0$
- What if that were at (2,2) location instead?

Our previous example

- We can use the following to change the order of the equations with $R_2 \leftrightarrow R_4$

```
A([2 4], :) = A([4 2], :); b([2 4]) = b([4 2]);
```

- Theoretically, the answer does not change, and the `\` gets it right
- But, `lufact` fails!

```
[L,U] = lufact(A);  
L
```

```
L =  
    1.0000         0         0         0  
   -1.0000    1.0000         0         0  
    0.5000        Inf    1.0000         0  
   -2.0000        Inf        NaN    1.0000
```

Our previous example

- We can use the following to change the order of the equations with $R_2 \leftrightarrow R_4$

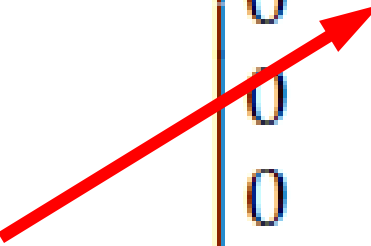
```
A([2 4], :) = A([4 2], :); b([2 4]) = b([4 2]);
```

- Theoretically, the answer does not change, and the `\` gets it right
- But, `lu` fails!

```
[L,U] = lu(A);  
L
```

```
L =  
    1.0000    0    0    0  
   -1.0000    1.0000    0    0  
    0.5000   Inf    1.0000    0  
   -2.0000   Inf   NaN    1.0000
```

- Why? Zero pivot after first column finished


$$\begin{bmatrix} 2 & 0 & 4 & 3 \\ 0 & 0 & 6 & -10 \\ 0 & 15 & 0 & -6 \\ 0 & 5 & 1 & -4 \end{bmatrix}$$

Fixing up issues with naïve GE

- The fact that we get a zero pivot can be fixed by switching rows, IF the column below the pivot is not all zero
- If the $A(i:n,i)=0$, then the columns $1:i$ are linearly dependent, and there is no unique solution
- This implies that the original matrix **A** is singular
- Theorem 2: If a pivot element and all the elements below it are zero, then the original matrix **A** is singular. In other words, if **A** is nonsingular, then Gaussian elimination with row pivoting will run to completion.
- This tells us that using row pivoting is well worth implementing

Our previous example

- How to swap rows?
- Use a “permutation matrix”
- If we take the identity matrix I that is 3x3, and switch the first two rows, we get

$$P = \begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

- The effect of left multiplying by this P is to switch rows 1 and 2! In terms

$$\mathbf{P}\mathbf{B} = \begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 2 & 3 & 1 \\ 6 & 5 & 4 \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 6 & 5 & 4 \\ 2 & 3 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$

- We could also look at \mathbf{P} as $\mathbf{P} = [\mathbf{e}_2 \quad \mathbf{e}_1 \quad \mathbf{e}_3]$

Permutations and row switches

- Permutation matrices have interesting properties
- Say left multiplying by P switches some rows; left multiplying by P^T will switch them back.
- This suggests that $P^T = P^{-1}$!! (Proof in exercises)
- Let's use them to keep track of row switches.
- For our example where we got the zero pivot, we had

$$\begin{bmatrix} 2 & 0 & 4 & 3 \\ 0 & 0 & 6 & -10 \\ 0 & 15 & 0 & -6 \\ 0 & 5 & 1 & -4 \end{bmatrix}$$

- We can switch rows 2 and 3 to keep going with GE
- Write this in matrix terms...

Row switches in LU

- After finishing with the first column and doing row switch, we have

$$PA_1 = P_1 L_{41} L_{31} L_{21} A$$

- Finishing up the factorization gives

$$U = L_{43} L_{42} L_{32} P_1 L_{41} L_{31} L_{21} A$$

- Working toward undoing the right side, we get

$$P_1^T L_{32}^{-1} L_{42}^{-1} L_{43}^{-1} U = L_{41} L_{31} L_{21} A$$

- Continuing, we get

$$L_{21}^{-1} L_{31}^{-1} L_{41}^{-1} P_1^T L_{32}^{-1} L_{42}^{-1} L_{43}^{-1} U = A$$

- The way that we use this is to post facto create P so that

$$LU = PA$$

- Or $P^T LU = A$

Row switches in LU: MATLAB

- We will use MATLAB's builtin function `lu`
- The syntax is $[L,U,P] = \text{lu}(A)$
- To use this, we first view the system at $Pax=Pb$
- Then use L and U as before, because row switching won't be needed now
- Thus, $PA = LU$ and let $Ux = z$.
- Then, $Lz = Pb$
- Solving the system can then be done by:
 - using `forwardsub` on $Lz=Pb$ to get z ,
 - then using `backsub` on $Ux=z$ to get x

Partial pivoting in MATLAB

- Solving systems using `lu` in MATLAB and our functions is then as follows:
 1. Find L , U and P from $[L,U,P] = \text{lu}(A)$;
 2. Solve $Lz=Pb$ using $z = \text{forwardsub}(L,P*b)$;
 3. Solve $Ux=z$ using $x = \text{backsub}(U,z)$.
- Using MATLAB's `lu` and `\`, one does:
 1. Find L , U and P from $[L,U,P] = \text{lu}(A)$;
 2. Solve $Lz=Pb$ using $z = L \setminus (P*b)$;
 3. Solve $Ux=z$ using $x = U \setminus z$.

Example using LU

```
A = [ 2 0 4 3; -2 0 2 -13 ; 1 15 2 -4.5 ; -4 5 -7 -10 ];  
b = [ 4; 40; 29; 9 ];
```

- Swapped row system at left

Example using LU

```
A = [ 2 0 4 3; -2 0 2 -13 ; 1 15 2 -4.5 ; -4 5 -7 -10 ];  
b = [ 4; 40; 29; 9 ];
```

```
[L,U,P] = lu(A)
```

```
L =  
    1.0000    0    0    0  
   -0.2500    1.0000    0    0  
    0.5000   -0.1538    1.0000    0  
   -0.5000    0.1538    0.0833    1.0000  
U =  
   -4.0000    5.0000   -7.0000  -10.0000  
         0   16.2500    0.2500   -7.0000  
         0         0    5.5385   -9.0769  
         0         0         0   -0.1667  
P =  
     0     0     0     1  
     0     0     1     0  
     0     1     0     0  
     1     0     0     0
```

- Swapped row system at left
- lu gives L,U,P



Example using LU

```
A = [ 2 0 4 3; -2 0 2 -13 ; 1 15 2 -4.5 ; -4 5 -7 -10 ];  
b = [ 4; 40; 29; 9 ];
```

```
[L,U,P] = lu(A)
```

```
L =  
    1.0000    0    0    0  
   -0.2500    1.0000    0    0  
    0.5000   -0.1538    1.0000    0  
   -0.5000    0.1538    0.0833    1.0000  
U =  
   -4.0000    5.0000   -7.0000  -10.0000  
         0   16.2500    0.2500   -7.0000  
         0         0    5.5385   -9.0769  
         0         0         0   -0.1667  
P =  
     0     0     0     1  
     0     0     1     0  
     0     1     0     0  
     1     0     0     0
```

- Swapped row system at left
- lu gives L,U,P
- In one line, forwardsub, then backsub

```
x = backsub( U, forwardsub( L, P*b) )
```

```
x =  
   -3.0000  
    1.0000  
    4.0000  
   -2.0000
```

Example using PtLU

```
[PtL,U] = lu(A)
```

```
PtL =  
  -0.5000    0.1538    0.0833    1.0000  
   0.5000   -0.1538    1.0000         0  
  -0.2500    1.0000         0         0  
   1.0000         0         0         0  
U =  
  -4.0000    5.0000   -7.0000  -10.0000  
         0   16.2500    0.2500   -7.0000  
         0         0    5.5385   -9.0769  
         0         0         0   -0.1667
```

```
x = U \ (PtL\b)
```

```
x =  
  -3.0000  
   1.0000  
   4.0000  
  -2.0000
```

- Use only two output arguments
- lu gives PtL and U
- The backslash is better than forwardsub and can use the row-switched L (i.e, PtL)
- Using the backslash can solve the two system sequence in one line

Example using only MATLAB's \

```
x = A\b
```

```
x =  
    -3  
     1  
     4  
    -2
```

```
x =  
-3.0000  
 1.0000  
 4.0000  
-2.0000
```

- Using \ solves the system directly
- Any row switching is transparent
- Compares well with previous solution

What about small pivots

- We have row swapped when the pivot was zero
- However, bad things happen when pivots are small ($\epsilon \ll 1$).
- The system at left has a small pivot at (1,1); do GE on the augmented matrix
- If we don't switch it out, then we get large multiplier and result at right
- If ϵ is small enough, then possible subtractive cancellation
- How to avoid?

$$\mathbf{A} = \begin{bmatrix} -\epsilon & 1 \\ 1 & -1 \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} 1 - \epsilon \\ 0 \end{bmatrix}$$

$$\begin{bmatrix} -\epsilon & 1 & 1 - \epsilon \\ 0 & -1 + \epsilon^{-1} & \epsilon^{-1} - 1 \end{bmatrix} \Rightarrow \begin{aligned} x_2 &= 1 \\ x_1 &= \frac{(1 - \epsilon) - 1}{-\epsilon} \end{aligned}$$

What about small pivots

$$\mathbf{A} = \begin{bmatrix} -\epsilon & 1 \\ 1 & -1 \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} 1 - \epsilon \\ 0 \end{bmatrix}$$

- Now try with rows swapped
- This time, multiplier is small, and we get a very good answer

$$\begin{bmatrix} 1 & -1 & 0 \\ 0 & 1 - \epsilon & 1 - \epsilon \end{bmatrix} \Rightarrow \begin{array}{l} x_2 = 1 \\ x_1 = \frac{0 - (-1)}{1} \end{array}$$

- No loss of significance this way
- This suggests a strategy: When working with row i , find the biggest element in column i from the diagonal and below, and put that element in the pivot
- That is $\max A(i:n,i)$ (say row p) becomes the pivot ($R_p \leftrightarrow R_i$)
- Then, multipliers are never bigger than 1, so this avoids the large multiplier problem

Vector Norms

- We need to be able to measure the size of a vector or matrix, so that we can order them
- We do this with norms, which can be thought of as functions that map $\mathbb{R}^n \mapsto \mathbb{R}$
- A norm of a vector $\|v\|$ must have the following properties:
 1. $\|x\| \geq 0$ for all $x \in \mathbb{R}^n$
 2. $\|x\| = 0$ if and only if $x = 0$
 3. $\|\alpha x\| = |\alpha| \|x\|$ for any $x \in \mathbb{R}^n$ and $\alpha \in \mathbb{R}$
 4. $\|x + y\| \leq \|x\| + \|y\|$ for all $x, y \in \mathbb{R}^n$ (the triangle inequality)

Norms of vectors

- The general vector norm of interest is the p -norm:

$$\|v\|_p = \left[\sum_{i=0}^n |v_i|^p \right]^{1/p}$$

- We are most interested in only three values of p : 1, 2, or ∞

$$\|\mathbf{x}\|_2 = \left(\sum_{i=1}^n |x_i|^2 \right)^{\frac{1}{2}} = \sqrt{\mathbf{x}^T \mathbf{x}}$$

$$\|\mathbf{x}\|_\infty = \max_{i=1, \dots, n} |x_i|$$

$$\|\mathbf{x}\|_1 = \sum_{i=1}^n |x_i|$$

Norms of vectors

- Examples:

$$u = [1 \ -5 \ 3], \quad v = \begin{bmatrix} -2 \\ 2 \\ 0 \end{bmatrix}$$

- Shape of vector doesn't matter
- 1-norm uses $p=1$, so that

$$\|u\|_1 = |1| + |-5| + |3| = 9, \quad \|v\|_1 = |-2| + |2| + 0 = 4$$

- For the 2-norm

$$\|u\|_2 = (|1|^2 + |-5|^2 + |3|^2)^{1/2} = \sqrt{35},$$

$$\|v\|_2 = (|-2|^2 + |2|^2 + 0)^{1/2} = \sqrt{8} = 2\sqrt{2}$$

Norms of vectors

- Examples:

$$u = [1 \ -5 \ 3], \quad v = \begin{bmatrix} -2 \\ 2 \\ 0 \end{bmatrix}$$


- ∞ -norm now:

$$\|u\|_{\infty} = \max(|1|, |-5|, |3|) = 5$$

- Multiply by a scalar?
- Always ≥ 0 ?
- Only 0 if **0** vector?

Norms of vectors

- What do unit vectors look like?
- Given $\mathbf{x} = [x_1 \ x_2]$, we can create the unit vector or direction $\mathbf{x}/\|\mathbf{x}\|$
- For $\mathbf{x} = [x_1 \ x_2]$ and $\|\mathbf{x}\|_2$, unit vectors have $\|\mathbf{x}\|_2 = (|x_1|^2 + |x_2|^2)^{1/2} = 1$
- In the (x_1, x_2) -plane, the set of all unit vector is a circle
- What about ∞ -norm?
- $\|\mathbf{x}\|_\infty = \max(|x_1|, |x_2|) = 1$; what is it?



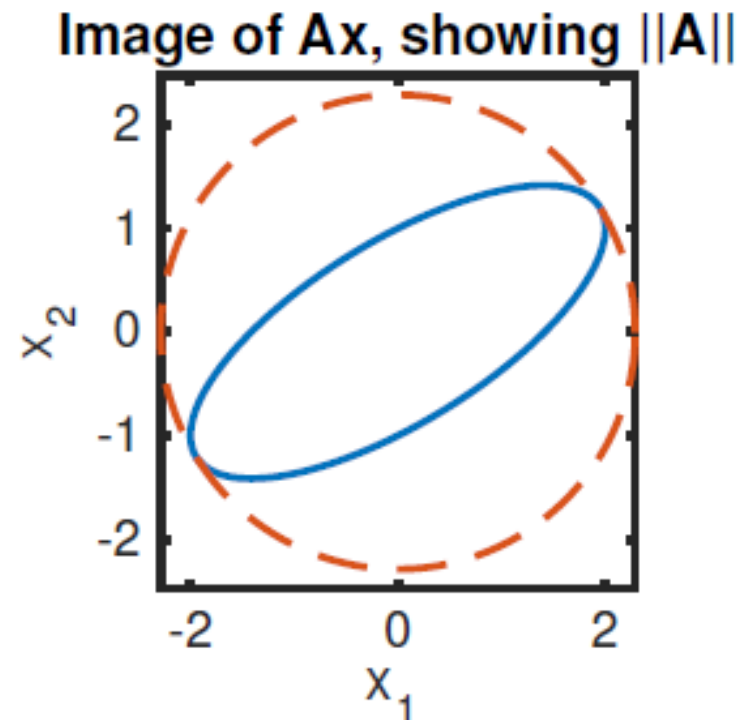

Norms of matrices

- Measuring the “size” of a matrix should be associated with a vector norm: Induced matrix norms
- For $\mathbf{x} = [x_1 \ x_2]$ and $\|\mathbf{x}\|_2$, the induced matrix norm is $\|\mathbf{A}\|_2 = \max_{\|\mathbf{x}\|_2=1} (\|\mathbf{A}\mathbf{x}\|_2) = \max_{\|\mathbf{x}\|_2 \neq 0} (\|\mathbf{A}\mathbf{x}\|_2 / \|\mathbf{x}\|_2)$
- When we multiply \mathbf{x} by a matrix, then it is stretched and pointed in a different direction, generally
- The magnitude of the biggest stretching is the norm of \mathbf{A}

Norms of matrices

- We can multiply the set of unit vectors by A and see what the distortion does to all of those unit vectors
- The magnitude of the biggest stretching is the norm of A

```
subplot(1,2,2), plot(Ax(1,:),Ax(2,:)), axis equal  
hold on, plot(twonorm*x(1,:),twonorm*x(2,:), '--')  
title('Image of Ax, showing ||A||')  
xlabel('x_1'), ylabel('x_2')
```



Norms of matrices

- The induced norms for the other matrices are a bit easier to compute
- The ∞ -norm is the maximum row sum of the matrix A
- The 1-norm is the maximum column sum of A
- Mnemonic: think of the “direction” of 1 or ∞
- Example: $A = \begin{bmatrix} 1 & 3 \\ -5 & 8 \end{bmatrix}$
- $\|A\|_1 = \max(1 + 5, 3 + 8) = 11$
- $\|A\|_\infty = \max(1 + 3, 5 + 8) = 13$

Norms of matrices

- We need some properties of matrix norms for future use.
- One can prove the following for any of the norms we have used.

For any $n \times n$ matrix \mathbf{A} and induced matrix norm,

$$\|\mathbf{Ax}\| \leq \|\mathbf{A}\| \cdot \|\mathbf{x}\|,$$

for any $\mathbf{x} \in \mathbb{R}^n$,

$$\|\mathbf{AB}\| \leq \|\mathbf{A}\| \cdot \|\mathbf{B}\|,$$

for any $\mathbf{B} \in \mathbb{R}^{n \times n}$,

$$\|\mathbf{A}^k\| \leq \|\mathbf{A}\|^k,$$

for any integer $k \geq 0$.

Stability analysis of solving systems

- We want to solve $A\mathbf{x} = \mathbf{b}$
- We want to analyze how robust our answers are to perturbations of A and \mathbf{b}
- To measure what happens to the sizes of vectors and matrices, use norms
- We know that $||\mathbf{b}|| = ||A\mathbf{x}|| \leq ||A|| ||\mathbf{x}||$
- We also know that $\mathbf{x} = A^{-1}\mathbf{b}$, now take norms
- $||\mathbf{x}|| = ||A^{-1}\mathbf{b}|| \leq ||A^{-1}|| ||\mathbf{b}||$
- If we change \mathbf{b} to $\mathbf{b} + \mathbf{d}$, then the solution changes somewhat to $\mathbf{x} + \mathbf{h}$, say, such that $A(\mathbf{x} + \mathbf{h}) = \mathbf{b} + \mathbf{d}$
- But, using the original equation $A\mathbf{x} = \mathbf{b}$, the first term on each side cancels

Stability analysis of solving systems

- We then have $\mathbf{A}\mathbf{h} = \mathbf{d}$
- Then $\mathbf{h} = \mathbf{A}^{-1}\mathbf{d}$, now take norms
- $\|\mathbf{h}\| = \|\mathbf{A}^{-1}\mathbf{d}\| \leq \|\mathbf{A}^{-1}\| \|\mathbf{d}\|$
- Let's now look at the relative change in the solution $\|\mathbf{h}\|/\|\mathbf{x}\|$ compared to the relative change in the right hand side $\|\mathbf{d}\|/\|\mathbf{b}\|$
- Taking the ratio of those two,

• The last quantity is the condition number

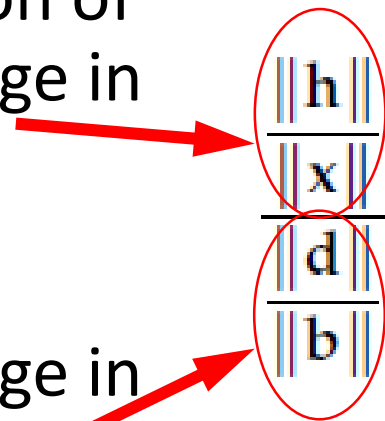
$$\kappa(\mathbf{A}) = \|\mathbf{A}^{-1}\| \|\mathbf{A}\|$$
$$\frac{\frac{\|\mathbf{h}\|}{\|\mathbf{x}\|}}{\frac{\|\mathbf{d}\|}{\|\mathbf{b}\|}} = \frac{\|\mathbf{h}\| \cdot \|\mathbf{b}\|}{\|\mathbf{d}\| \cdot \|\mathbf{x}\|} \leq \frac{(\|\mathbf{A}^{-1}\| \|\mathbf{d}\|)(\|\mathbf{A}\| \|\mathbf{x}\|)}{\|\mathbf{d}\| \|\mathbf{x}\|} = \|\mathbf{A}^{-1}\| \|\mathbf{A}\|$$

Stability analysis of solving systems

The condition number
 $\kappa(\mathbf{A}) = \|\mathbf{A}^{-1}\| \|\mathbf{A}\|$
thus tells us the worst
case magnification of
the relative change in
the answer

compared to

the relative change in
the right hand side


$$\frac{\frac{\|\mathbf{h}\|}{\|\mathbf{x}\|}}{\frac{\|\mathbf{d}\|}{\|\mathbf{b}\|}} = \frac{\|\mathbf{h}\| \cdot \|\mathbf{b}\|}{\|\mathbf{d}\| \cdot \|\mathbf{x}\|} \leq \frac{(\|\mathbf{A}^{-1}\| \|\mathbf{d}\|) (\|\mathbf{A}\| \|\mathbf{x}\|)}{\|\mathbf{d}\| \|\mathbf{x}\|} = \|\mathbf{A}^{-1}\| \|\mathbf{A}\|$$

The condition number depends on the matrix \mathbf{A} and the norm, and it measures the sensitivity of its solutions to perturbations

Stability analysis of solving systems

- What if we perturb A and solve $Ax = b$?
- If we change A to $A + E$, then the solution changes somewhat to $x + h$, say, such that $(A + E)(x + h) = b$
- Expand again to get $Ax + Ex + Ah + Eh = b$
- Then use $Ax = b$, and neglect Eh because it is the product of two small terms; then one has $Ex + Ah = 0$
- Solve for h and use norms:

$$h = -A^{-1}Ex$$

$$\|h\| \leq \|A^{-1}\| \|E\| \|x\|$$

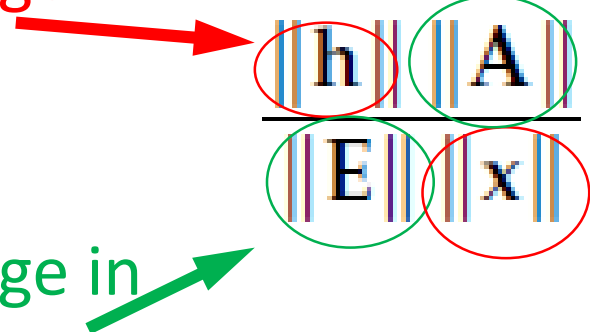
- Now compute relative changes again...

Stability analysis of solving systems

The condition number
 $\kappa(\mathbf{A}) = \|\mathbf{A}^{-1}\| \|\mathbf{A}\|$
thus tells us the worst
case magnification of
the **relative change in
the answer**

compared to

the **relative change in
the matrix \mathbf{A}**


$$\frac{\frac{\|\mathbf{h}\|}{\|\mathbf{E}\|}}{\frac{\|\mathbf{A}\|}{\|\mathbf{x}\|}} \leq \frac{\left(\|\mathbf{A}^{-1}\| \|\mathbf{E}\| \|\mathbf{x}\| \right) \|\mathbf{A}\|}{\|\mathbf{E}\| \|\mathbf{x}\|} = \kappa(\mathbf{A}).$$

The condition number measures the sensitivity of its solutions to perturbations in the matrix as well

Condition number

- In summary, the worst case changes in the answer from perturbations to the coefficients is given by the condition number:

$$\frac{\|\Delta \mathbf{x}\|}{\|\mathbf{x}\|} \leq \kappa(\mathbf{A}) \frac{\|\Delta \mathbf{b}\|}{\|\mathbf{b}\|}$$
$$\frac{\|\Delta \mathbf{x}\|}{\|\mathbf{x}\|} \leq \kappa(\mathbf{A}) \frac{\|\Delta \mathbf{A}\|}{\|\mathbf{A}\|}$$

- Note that the best we can get is a condition number of unity:

$$1 = \|\mathbf{I}\| = \|\mathbf{A}\mathbf{A}^{-1}\| \leq \|\mathbf{A}\| \|\mathbf{A}^{-1}\| = \kappa(\mathbf{A})$$

Condition number: example

- The $n \times n$ Hilbert matrix has elements

$$H_{ij} = 1/(i + j - 1)$$

- It's a builtin function in MATLAB:

```
>> hilb(5)
ans =
    1.0000    0.5000    0.3333    0.2500    0.2000
    0.5000    0.3333    0.2500    0.2000    0.1667
    0.3333    0.2500    0.2000    0.1667    0.1429
    0.2500    0.2000    0.1667    0.1429    0.1250
    0.2000    0.1667    0.1429    0.1250    0.1111
```

- It is not singular, but conditioning?

```
>> det(hilb(5))
ans =
    3.7493e-12
>> cond(hilb(5))
ans =
    4.7661e+05
```

Condition number: example

- Explore the Hilbert matrix for solving systems
- Increase n gets even worse conditioning fast.
- The “gallery” has a collection of commonly used matrices in numerical methods and analysis.
- Try “help gallery” at prompt or search for gallery in help browser
- You can find other ill-conditioned, non-singular matrices there: e.g., dorr,

Residuals and ill-conditioned systems

- What happens to the residual $\mathbf{r} = \mathbf{b} - \mathbf{A}\mathbf{x}$?
- For the computed solution (tilde) we have $\mathbf{r} = \mathbf{b} - \mathbf{A}\tilde{\mathbf{x}}$.
- Manipulate into the useful form
$$\mathbf{r} = \mathbf{A}(\mathbf{A}^{-1}\mathbf{b} - \tilde{\mathbf{x}}) = \mathbf{A}(\mathbf{x} - \tilde{\mathbf{x}}),$$
- Then use solve for change in solution and use norms: $\|\mathbf{x} - \tilde{\mathbf{x}}\| \leq \|\mathbf{A}^{-1}\| \|\mathbf{r}\|.$
- Now compute relative changes again:
$$\frac{\|\mathbf{x} - \tilde{\mathbf{x}}\|}{\|\mathbf{x}\|} \leq \kappa(\mathbf{A}) \frac{\|\mathbf{r}\|}{\|\mathbf{b}\|}$$
- We can expect *small residual*, but...
the *change in the answer may not be small* for ill-conditioned systems!!!
- Commercial FEM, e.g., tells you *residuals*: use with care!!!

Special matrices

- We can take advantage of the structure of some kinds of matrices to create faster algorithms or get additional information.
- Possibilities include:
 - Triangular matrices: Been there, reduces solves for n unknowns from $O(n^3)$ to $O(n^2)$
 - Banded matrices: only certain diagonals have non-zero elements
 - Sparse matrices: only a (preferably small) minority of elements are nonzero
 - Symmetric matrices
 - Symmetric positive definite matrices

Banded matrices

- A tridiagonal matrix has three nonzero diagonals: e.g.,

```
A = gallery('tridiag',c,d,e)
```

- This matrix would have vectors c in the first subdiagonal, d on the main diagonal, and e on the first superdiagonal
- The bandwidth is three. How to compute?
- Upper bandwidth is p , with $A_{ij} = 0$ if $j - i > p$
- Lower bandwidth is q , with $A_{ij} = 0$ if $i - j > q$
- Total bandwidth is $p + q - 1$
- For this example, we'd get 3, with $p = 2$ and $q = 2$

Banded matrices

- Consider

```
>> triu(tril(rand(7),1),-3)
```


```
ans =
```

0.4465	0.3725	0	0	0	0	0
0.6463	0.5932	0.4067	0	0	0	0
0.5212	0.8726	0.6669	0.9880	0	0	0
0.3723	0.9335	0.9337	0.8641	0.5583	0	0
0	0.6685	0.8110	0.3889	0.5989	0.8825	0
0	0	0.4845	0.4547	0.1489	0.2850	0.2834
0	0	0	0.2467	0.8997	0.6732	0.8962


- Upper bandwidth is $p = 1$, lower bandwidth is $q = 3$
- Total bandwidth is $p + q + 1 = 5$

Banded matrices

- We can visualize nonzero elements with sparsity plot from `spy(triu(tril(rand(7),1,-3))`
- Dots are shown where the nonzero elements were
- The total number of nonzeros is at the bottom
- Handy to see where nonzeros are in large matrices



Banded matrices

- How do we get diagonals?
- How can we build banded matrices
- What happens with LU factorization?
- Start with example at right
- We can extract one diagonal at a time with the diag command 

```
A =  
    2    -1     0     0     0     0  
    4     2    -1     0     0     0  
    0     3     0    -1     0     0  
    0     0     2     2    -1     0  
    0     0     0     1     1    -1  
    0     0     0     0     0     2
```

```
diag_main = diag(A,0)'
```

```
diag_main =  
    2     2     0     2     1     2
```

```
diag_plusone = diag(A,1)'
```

```
diag_plusone =  
   -1    -1    -1    -1    -1
```

```
diag_minusone = diag(A,-1)'
```


```
diag_minusone =  
    4     3     2     1     0
```

Banded matrices

- How do we get diagonals?
- How can we build banded matrices?
- What happens with LU factorization?
- Start with example at right
- We can extract one diagonal at a time with the diag command
- We can also create the matrix with the diag command:

A =

2	-1	0	0	0	0
4	2	-1	0	0	0
0	3	0	-1	0	0
0	0	2	2	-1	0
0	0	0	1	1	-1
0	0	0	0	0	2



```
>> c = [4 3 2 1 0];  
>> d = [2 2 0 2 1 2];  
>> e = [-1 -1 -1 -1 -1];  
>> B = diag(c,-1)+diag(d,0)+diag(e,1);  
>> norm(A-B,2)  
ans =  
0
```

Banded matrices



- What happens with LU factorization?
- Modify A with extra diag
- Do LU factorization without partial pivoting
- The banded structure is preserved:

```
A = A + diag([5 8 6 7],2)
```

```
A =
```

2	-1	5	0	0	0
4	2	-1	8	0	0
0	3	0	-1	6	0
0	0	2	2	-1	7
0	0	0	1	1	-1
0	0	0	0	0	2

```
[L,U] = lufact(A);  
subplot(1,2,1), spy(L), title('L')  
subplot(1,2,2), spy(U), title('U')
```



Banded matrices



- What happens with LU factorization?
- Modify A with extra diag
- Now do LU factorization with partial pivoting
- The banded structure is not preserved:
- We can improve the computing time by telling matlab that the matrices may be sparse (mostly zeros)

```
A = A + diag([5 8 6 7],2)
```

```
A =
```

2	-1	5	0	0	0
4	2	-1	8	0	0
0	3	0	-1	6	0
0	0	2	2	-1	7
0	0	0	1	1	-1
0	0	0	0	0	2

```
[L,U,P] = lu(A);  
subplot(1,2,1), spy(L), title('L')  
subplot(1,2,2), spy(U), title('U')
```




Sparse matrices

- Consider tridiagonal matrices

```
A = gallery('tridiag',n)
```

- This creates a $n \times n$ tridiagonal matrix with 2 on the main diagonal and -1 on the sub- and super-diagonals
- But, only the non-zeros are stored (let $n=10$ and try it!) Thus, only 28 numbers are stored for this element that would have 100 elements: 28%
- If $n = 100$, then $\text{nz}=298$, out of 10^4 possible elements: 2.98% nonzero
- In this case, $3n - 2$ elements nonzero out of n^2 possible; more sparse as n increases

```
>> A = gallery('tridiag',10);  
>> spy(A)
```



Sparse matrices

- Computations can be sped up significantly if we only work with the nonzeros (need to code it)

```
n = 8000;  
A = diag(1:n) + diag(n-1:-1:1,1) + diag(ones(n-1,1),-1);  
  
tic, [L,U] = lu(A); toc
```

```
Elapsed time is 5.850532 seconds.
```

```
tic, [L,U] = lu(sparse(A)); toc
```

```
Elapsed time is 0.182472 seconds.
```


- If A is sparse, full(A) will make it a full matrix

Sparse matrices

- Consider tridiagonal matrices

`A = gallery('wathen', n, n)`



- This creates a large sparse matrix; if $n = 20$, the matrix is 1281×1281 with
- But, things can go wrong: LU factorization leads to fill-in
- In this example, about 1.17% nonzeros in A , but look how many nonzeros after LU...



Sparse matrices

- The commands:
- Each factor L (left) and U (right) has about $6.6e4$ nonzeros now! (10x more nonzeros)

```
>> A=gallery('wathen',20,20);  
>> spy(A)  
>> [L,U]=lu(A);  
>> spy(L)  
>> spy(U)
```



Symmetric matrices

- If $\mathbf{A}^T = \mathbf{A}$, then \mathbf{A} is symmetric

- We can modify the LU factorization:

$$\mathbf{A} = \mathbf{L}\mathbf{U} = \mathbf{L}\mathbf{I}\mathbf{U} = \mathbf{L}\mathbf{D}\mathbf{D}^{-1}\mathbf{U}$$

- \mathbf{D} is diagonal, with the elements that would have been in \mathbf{U} from standard factorization
- Now it turns out that $\mathbf{D}^{-1}\mathbf{U} = \mathbf{L}^T$
- Then $\mathbf{A} = \mathbf{L}\mathbf{D}\mathbf{L}^T$

Symmetric positive definite matrices

- If $\mathbf{A}^T = \mathbf{A}$, and
- if $\mathbf{x}^T \mathbf{A} \mathbf{x} \geq 0$ for all $\mathbf{x} \in \mathbb{R}^n$ and $\mathbf{x} \neq 0$,
- Then \mathbf{A} is positive definite
- We can modify the LU factorization again to the Cholesky factorization:

$$\mathbf{A} = \mathbf{R}^T \mathbf{R}, \quad \mathbf{R} = \mathbf{D}^{1/2} \mathbf{L}^T$$

- \mathbf{R} has all positive entries on the diagonal
- The MATLAB function `chol` will compute the Cholesky factorization

Symmetric positive definite matrices

- Start with `A=magic(5)` ;

```
B = A'*A
```

```
B =
```

1055	865	695	770	840
865	1105	815	670	770
695	815	1205	815	695
770	670	815	1105	865
840	770	695	865	1055

```
R = chol(B)
```

```
R =
```

32.4808	26.6311	21.3973	23.7063	25.8615
0	19.8943	12.3234	1.9439	4.0856
0	0	24.3985	11.6316	3.7415
0	0	0	20.0982	9.9739
0	0	0	0	16.0005

```
norm( R'*R - B )
```

```
ans =
```

```
0
```

Symmetric positive definite matrices

- If $\mathbf{A}^T = \mathbf{A}$, and
- if $\mathbf{x}^T \mathbf{A} \mathbf{x} \geq 0$ for all $\mathbf{x} \in \mathbb{R}^n$ and $\mathbf{x} \neq 0$,
- Then \mathbf{A} is positive definite
- We can modify the LU factorization again to the Cholesky factorization:

$$\mathbf{A} = \mathbf{R}^T \mathbf{R}, \quad \mathbf{R} = \mathbf{D}^{1/2} \mathbf{L}^T$$

- \mathbf{R} has all positive entries on the diagonal
- The MATLAB function `chol` will compute the Cholesky factorization